

# Deploying full prometheus stacks via juju charms, *lightning'ly*

~ JuanJo Ciarlante - aka *jjo* ~

[@xjjo](#)

[jjo@canonical.com](mailto:jjo@canonical.com)

~ [promcon.io](http://promcon.io) - 2016-08-25 ~

# About me

- Working at Canonical Ltd as Cloud Reliability Engineer
  - since 2012
- Worked at Google Switzerland as Site Reliability Engineer
  - between 2007-2012
- FOSS contributor
  - since 1995

*So ... you don't like the ':' used for linux interfaces aliases ?*  
=> blame the guy at stage

# Juju in a few words

- service orchestration tool
- multi-provider:
  - clouds: AWS, GCE, Azure, Rackspace, Openstack
  - metals: MaaS
  - containers: LXC or LXD (on v2)
- services get deployed and setup by charms:
  - set of scripts invoked at different stages

# Juju orchestration - webapp deploy

juju  
controller

```
$ juju bootstrap
```

# Juju orchestration - webapp deploy



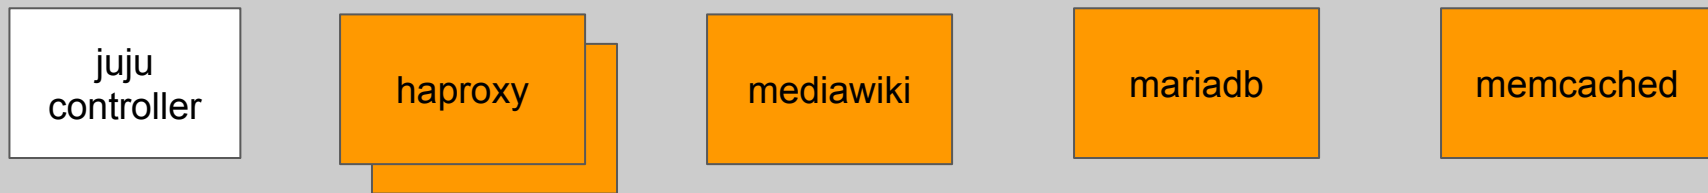
juju  
controller

haproxy

```
$ juju bootstrap
```

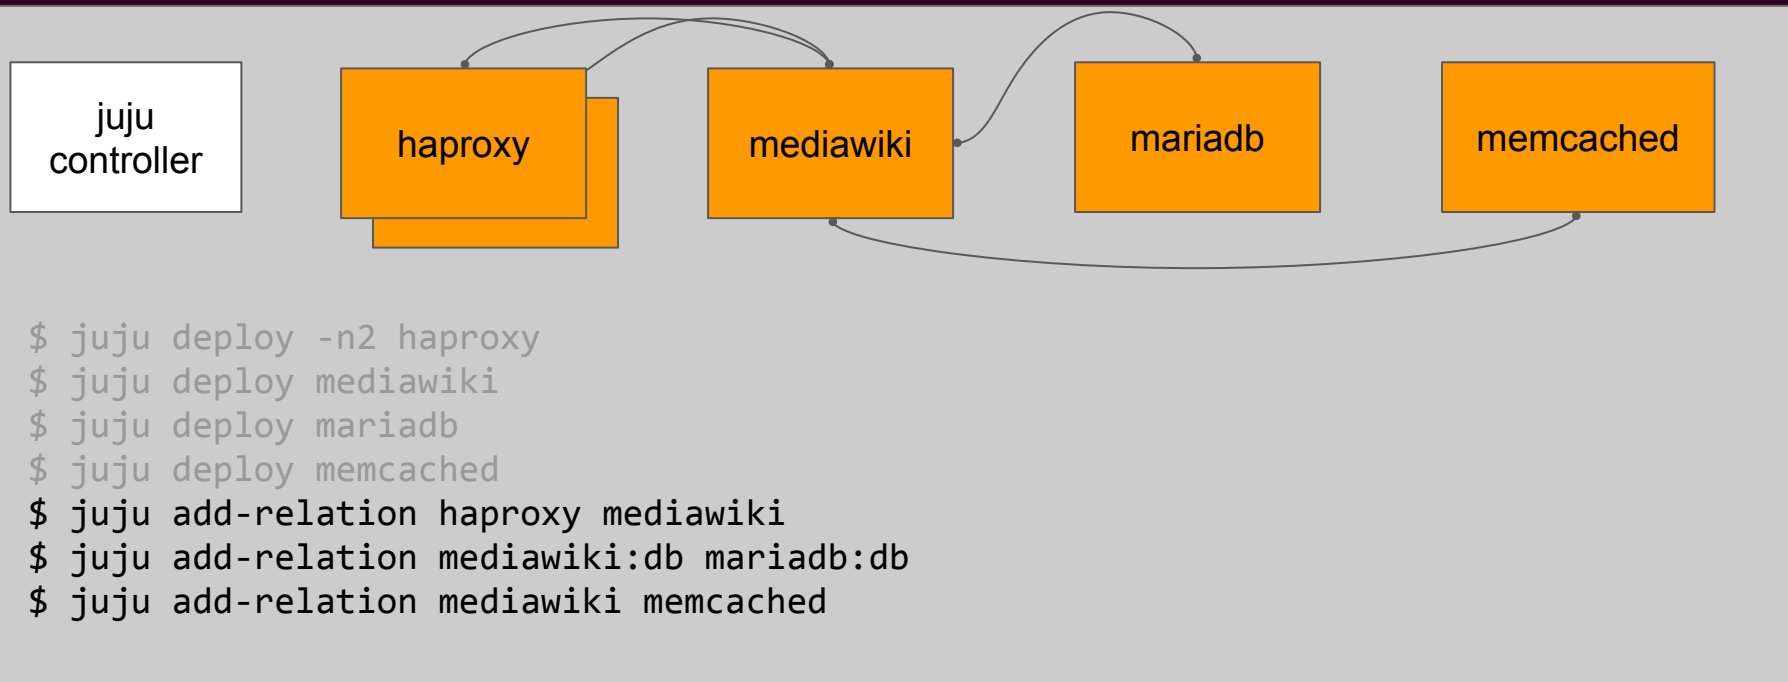
```
$ juju deploy -n2 haproxy
```

# Juju orchestration - webapp deploy

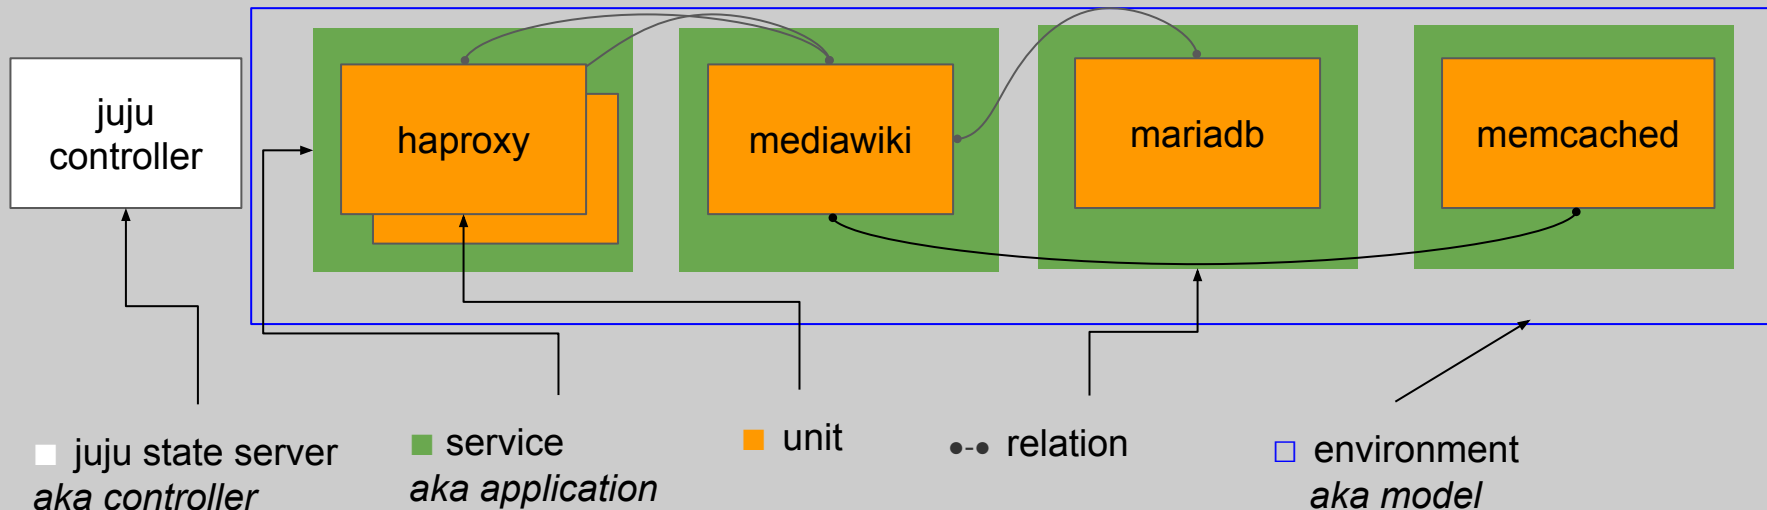


```
$ juju bootstrap
$ juju deploy -n2 haproxy
$ juju deploy mediawiki
$ juju deploy mariadb
$ juju deploy memcached
```

# Juju orchestration - webapp deploy



# Juju orchestration - glossary~ish



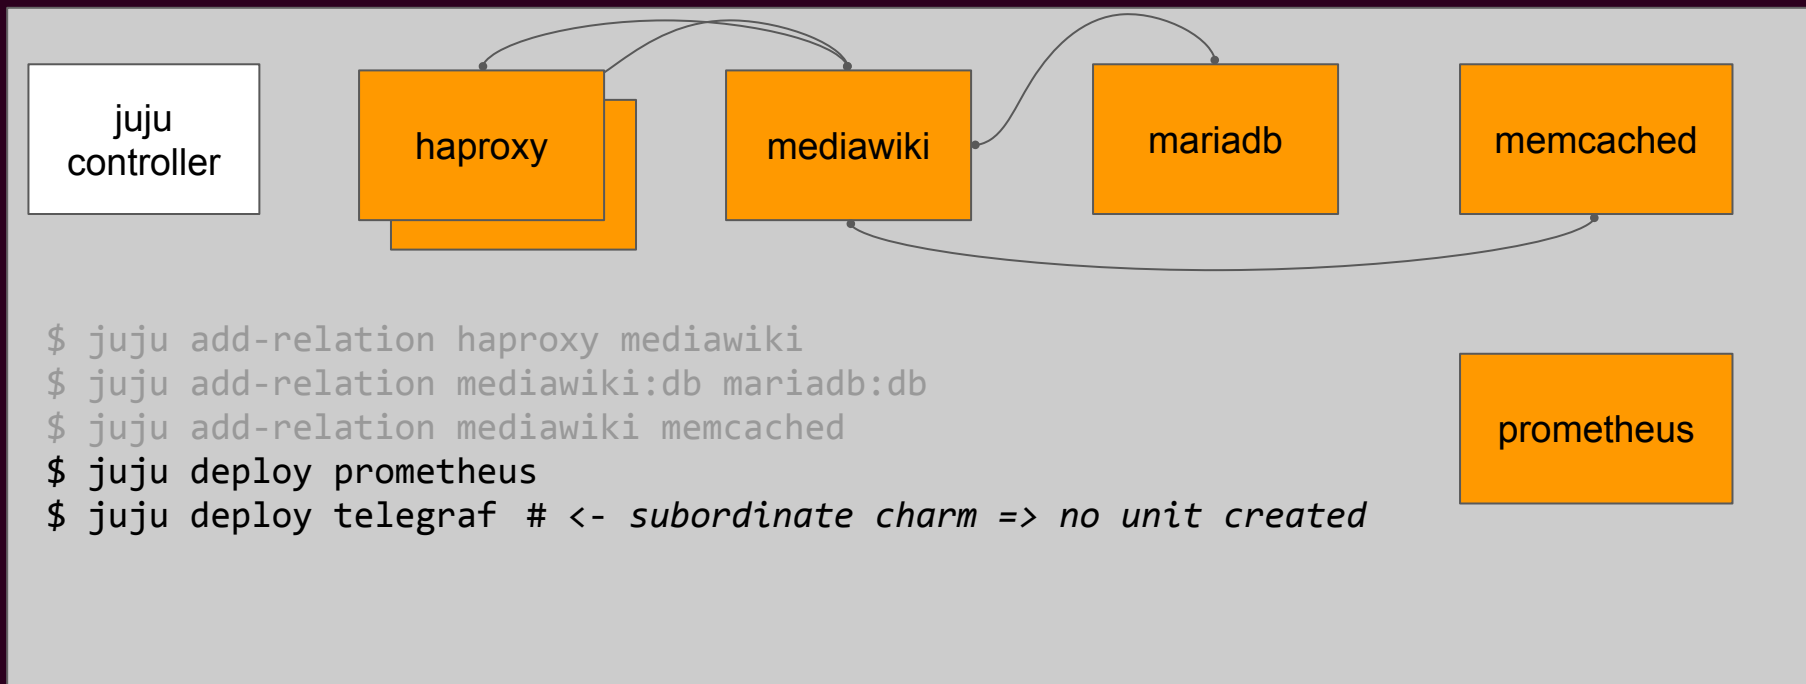
- **charm** code is specified by URL, then locally **run** at **units**
- service stacks can be specified as a **juju bundle** file (e.g. wiki-bundle.yaml)



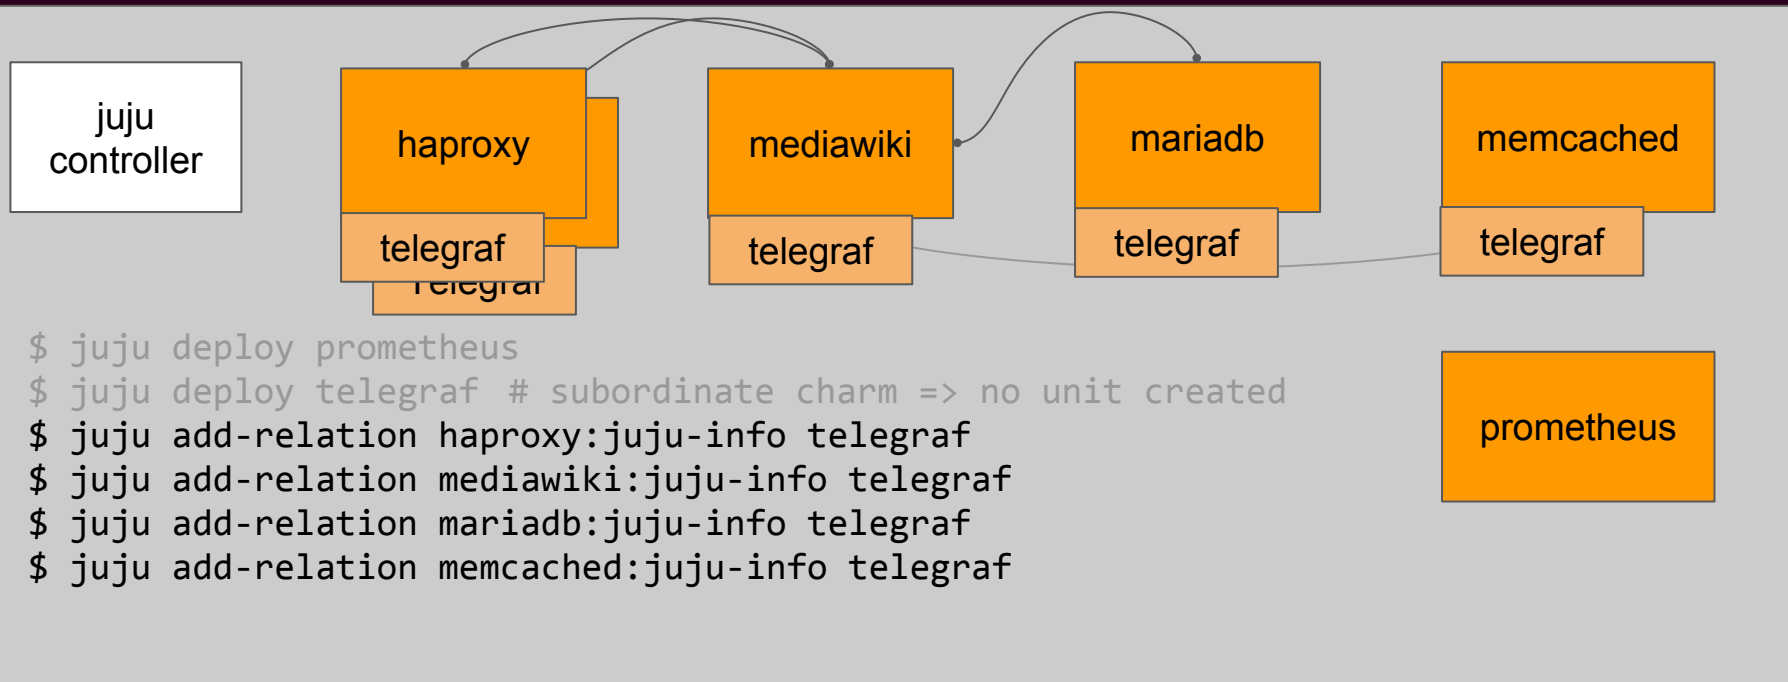
# What we use prometheus for

- **Infrastructure (phy nodes and overcloud services)**
  - /metrics: snmp exp., pushgateway, collectd exp., telegraf
  - maint. by: ops
- **Openstack (and storage-layer) metrics**
  - /metrics: collectd exp., pushgateway, telegraf, (wip) mtail
  - maint. by: ops
- **Online services**
  - /metrics: telegraf
  - maint. by: devs, ops

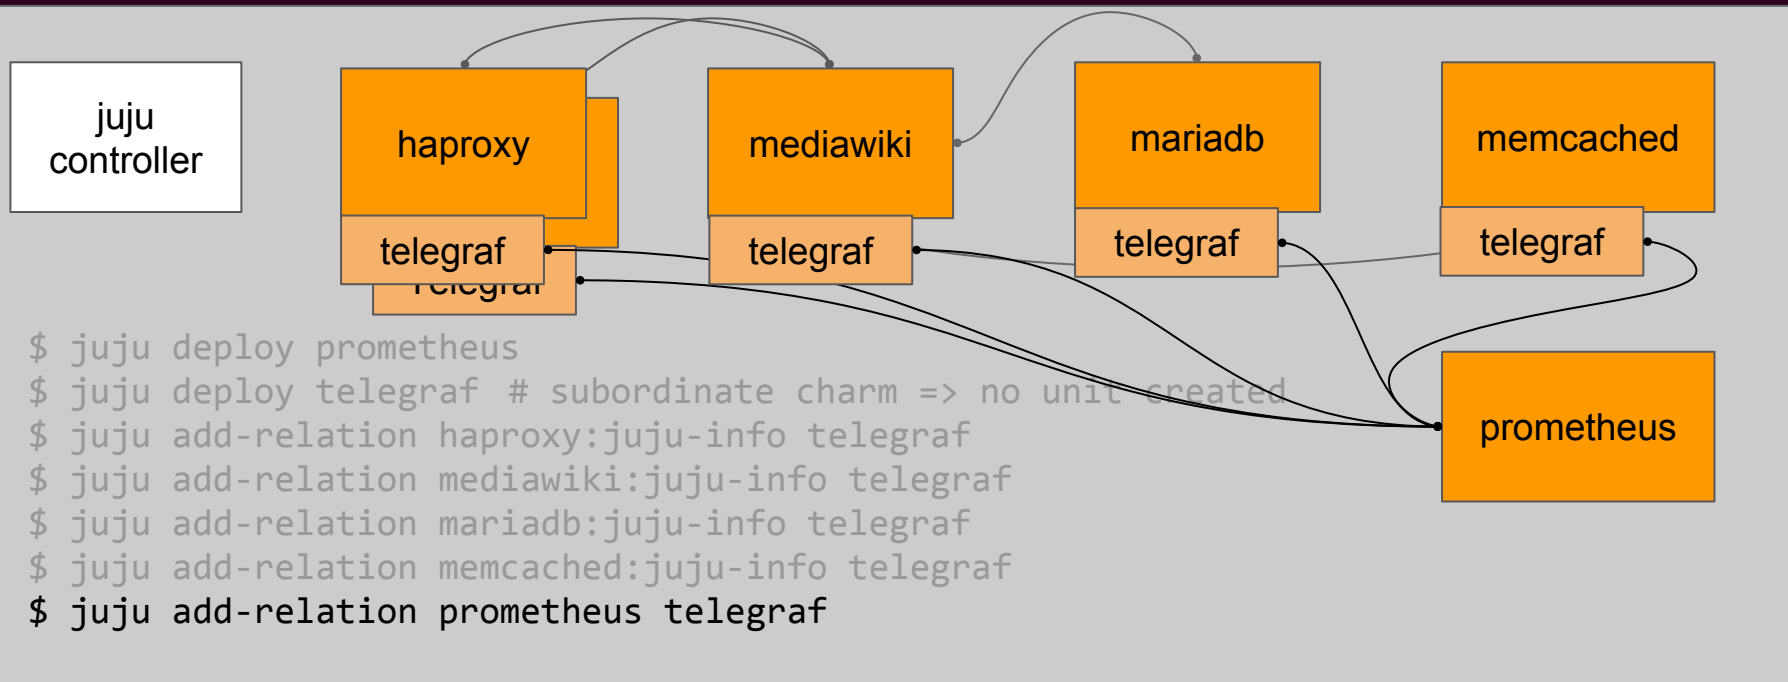
# How we deploy prometheus



# How we deploy prometheus

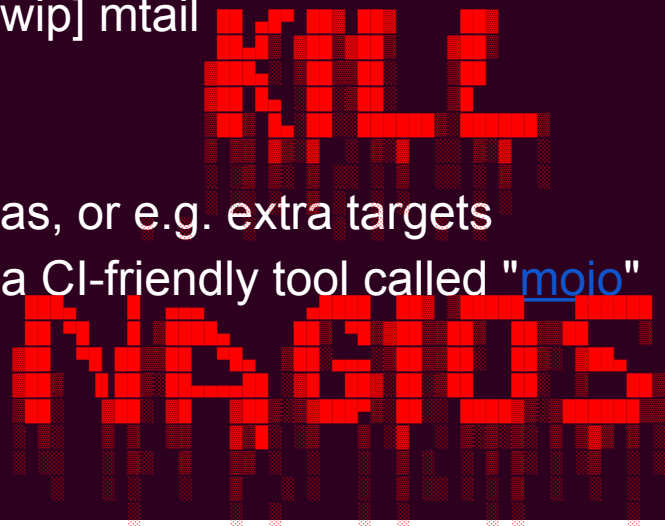


# How we deploy prometheus



# How we deploy prometheus (wrap-up)

- all these charms and required deb packages are publicly available as FOSS
- charms we developed, and use:
  - prometheus
  - telegraf, snmp-exporter, collectd-exporter, [wip] mtail
  - grafana
  - [wip] alertmanager
- most charms allow specifying extra config stanzas, or e.g. extra targets
- our deploys are specified by *bundles*, driven by a CI-friendly tool called "[moio](#)"
- more details w/ examples, next week at:
  - <http://github.com/jjo/promcon>



top: 20

service: .\*nova.\*

host: .\*

disk: sd[b-z][0-9]\*

iface: [eb].\*

rate\_type: irate

rate\_min\_bps: 10e3


rate\_min\_pps: 100

lat\_min\_ms: 1

LOAD PERFORMANCE METRICS


NETWORK PERFORMANCE METRICS [IRATE]

.\*nova.\* (hosts: .\*): network rx irate by host-iface (top 20 thru time) > 10e3



	max	avg	current
eth0	98.2 MBps	24.7 MBps	5.3 MBps
eth1	97.3 MBps	22.5 MBps	4.4 MBps
eth2	89.2 MBps	9.9 MBps	7.6 MBps
eth3	51.8 MBps	13.0 MBps	24.5 MBps
eth4	51.0 MBps	17.6 MBps	33.1 MBps
eth5	49.9 MBps	17.2 MBps	32.7 MBps
eth6	47.9 MBps	9.7 MBps	4.6 MBps
eth7	23.7 MBps	6.1 MBps	2.6 MBps
eth8	23.1 MBps	5.5 MBps	8.3 MBps

.\*nova.\* (hosts: .\*): network tx irate by host-iface (top 20 thru time) > 10e3



	max	avg	current
eth0	54.6 MBps	14.7 MBps	26.5 MBps
eth1	54.4 MBps	14.6 MBps	26.4 MBps
eth2	53.1 MBps	5.4 MBps	4.2 MBps
eth3	48.4 MBps	15.9 MBps	30.3 MBps
eth4	43.0 MBps	4.2 MBps	3.7 MBps
eth5	33.4 MBps	5.3 MBps	2.2 MBps
eth6	29.6 MBps	10.7 MBps	5.3 MBps
eth7	23.3 MBps	3.4 MBps	3.8 MBps
eth8	23.0 MBps	7.9 MBps	1.7 MBps

DISKS PERFORMANCE METRICS - THROUGHPUT IOPS [IRATE]

.\*nova.\* (hosts: .\*): read IOPs irate (top 20 over time) > 100

&& thank you :)  
@xjjo